

Example 10:

cqlsh

**Restrictions on partition
key**

Let's revisit the column definition of product

```
cassandra@cqlsh:catalog> CREATE COLUMNFAMILY product  
    productId varchar,  
    title text,  
    brand varchar,  
    publisher varchar,  
    length int,  
    breadth int,  
    height int,  
    PRIMARY KEY(productId)  
);
```

**Here productId is our partition key
and we have not defined any
clustering columns**

We want to list the books published by Riverhead Books


```
cassandra@cqlsh:catalog> select productid from product where publisher='Riverhead';
```

```
InvalidRequest: code=2200 [Invalid query] message="Cannot execute this query as it might involve data filtering and thus may have unpredictable performance. If you want to execute this query despite the performance unpredictability, use ALLOW FILTERING"
```

What went wrong?


The Murmur3Partitioner distributes data uniformly
across the cluster

So products of the book category are scattered
across the cluster



```
cassandra@cqlsh:catalog> select productid from product where publisher='Riverhead';
```

For a query like this, the coordinator node will send requests to **ALL** the nodes in the cluster to get the result



And here if **even 1** intermediate request fails, then the client request results in an **error**


```
cassandra@cqlsh:catalog> select productid from product where publisher='Riverhead';
```

if this query was run, cassandra would have to scan the
entire column family across **all** nodes

**Read will timeout while executing this
query**

To prevent such inefficient queries

cassandra has restrictions on

- the columns which can be queried
- the operations that are supported by the columns

Let's go through the restrictions on
the partition key

Partition key restrictions

Let's create a columnfamily "skus" which stores metadata about the products

```
cassandra@cqlsh:catalog> CREATE COLUMNFAMILY skus (  
    sellerId varchar,  
    productId varchar,  
    skuId varchar,  
    title text,  
    listingId varchar,  
    isListingCreated boolean,  
    timeofskucreation text,  
    PRIMARY KEY((sellerId,skuId), timeofskucreation, productId));
```


Partition key restrictions

```
cassandra@cqlsh:catalog> CREATE COLUMNFAMILY skus (  
    sellerId varchar,  
    productId varchar,  
    skuId varchar,  
    title text,  
    listingId varchar,  
    isListingCreated boolean,  
    timeofskucreation text,  
    PRIMARY KEY (sellerId,skuId), timeofskucreation, productId));
```

(sellerId,skuId) make the
partition key

Partition key restrictions

```
cassandra@cqlsh:catalog> CREATE COLUMNFAMILY skus (  
    sellerId varchar,  
    productId varchar,  
    skuId varchar,  
    title text,  
    listingId varchar,  
    isListingCreated boolean,  
    timeofskucreation text,  
    PRIMARY KEY((sellerId,skuId), timeofskucreation, productId));
```

**timeofskucreation and productId are the
clustering columns**

Partition key restrictions

We have data for some SKUs in a csv file

```
Next,COM1,SKU1,Acer One,LISTINGNextCOM1,true,2015-08-12 12:21
Next,COM2,SKU2,Acer One,LISTINGNextCOM2,true,2015-03-01 11:11
Chroma,COM1,CHROMASKU1,Acer One,,false,2016-01-01 00:01
Fab,S0FA1,FABSKU,Urban Living Derby,LISTINGFabS0FA1,true,2015-12-11 12:21
Decor,S0FA1,DECORSKU,Urban Living Derby,,false,2016-07-01 22:30
S0FA3,Urban Living 4 seater,Fab,FABSKU,2015-12-11 12:21,False
S0FA2,Urban Living 2 seater,Fab,FABSKU1,2015-12-11 12:21,False
```

Partition key restrictions

We can import this data in the columnfamily(CF) by using the **COPY** command

```
cassandra@cqlsh:catalog> COPY skus (sellerid, productid, skuid, title, listingid, islistingcreated, timeofskucreation) FROM 'SKUS.csv' ;
```

We provide the columns for which the file has data

Partition key restrictions

```
cassandra@cqlsh:catalog> select * from skus;
```

sellerid id	skuid title	timeofskucreation	productid	islistingcreated	listing
Chroma null	CHROMASKU1 Acer One	2016-01-01 00:01	COM1	False	
Decor null	DECORSKU Urban Living Derby	2016-07-01 22:30	SOFA1	False	
Next NextCOM2	SKU2 Acer One	2015-03-01 11:11	COM2	True	LISTING
Fab null	FABSKU1 Urban Living 2 seater	2015-12-11 12:21	SOFA2	False	
Next NextCOM1	SKU1 Acer One	2015-08-12 12:21	COM1	True	LISTING
Fab FabSOFA1	FABSKU Urban Living Derby	2015-12-11 12:21	SOFA1	True	LISTING
Fab null	FABSKU Urban Living 4 seater	2015-12-11 12:21	SOFA3	False	

(7 rows)

data is imported in the skus CF

Partition key restrictions

Let's query with only 1 column of the partition key - sellerid

```
cassandra@cqlsh:catalog> select * from skus where sellerid = ('Decor');
```

```
InvalidRequest: code=2200 [Invalid query] message="Partition key parts: skuid must be restricted as other parts are"
```

What went wrong?

Partition key restrictions

Let's query with only 1 column of the partition key - sellerid

```
cassandra@cqlsh:catalog> select * from skus where sellerid = ('Decor');
```

```
InvalidRequest: code=2200 [Invalid query] message="Partition key parts: skuid must be restricted as other parts are"
```

token for the row is generated by hashing the data
of the partition key

For the skus column family, data of the partition key is
⟨sellerid data, skuid data⟩

Partition key restrictions

Let's query with only 1 column of the partition key - sellerid

```
cassandra@cqlsh:catalog> select * from skus where sellerid = ('Decor');
```

output

```
InvalidRequest: code=2200 [Invalid query] message="Partition key parts: skuid must be restricted as other parts are"
```

Without skuid data, the token cannot be generated

And without a token, coordinator node will not be able to know which node owns the partition

Partition key restrictions

Let's query with only 1 column of the partition key - sellerid

```
cassandra@cqlsh:catalog> select * from skus where sellerid = ('Decor');
```

output

```
InvalidRequest: code=2200 [Invalid query] message="Partition key parts: skuid must be restricted as other parts are"
```

Let's understand the error message

Partition key restrictions

Let's query with only 1 column of the partition key - sellerid

```
cassandra@cqlsh:catalog> select * from skus where sellerid = ('Decor');
```

output

```
InvalidRequest: code=2200 [Invalid query] message="Partition key parts: skuid must  
be restricted as other parts are"
```

restricted here means that
definite data should be provided
for the column in the query

Query condition is required
on the skuid column as well
i.e the skuid column should be
restricted

Partition key restrictions

Let's query with only the sellerid

```
cassandra@cqlsh:catalog> select * from skus where sellerid = ('Decor');
```

output

InvalidRequest: code=2200 [Invalid query] message="Partition key parts: skuid must be restricted as other parts are"

**All the columns of the partition key
should be restricted in the query**

Partition key restrictions

All the columns of the partition key
should be restricted in the query

Partition key restrictions

Similarly IN query will fail, if all the columns of partition key are not restricted

```
cassandra@cqlsh:catalog> select * from skus where sellerid IN ('Decor');
```

```
InvalidRequest: code=2200 [Invalid query] message="Partition key parts: skuid must be restricted as other parts are"
```

Partition key restrictions

Let's do a range query using '>' on the partition key

```
cassandra@cqlsh:catalog> select * from skus where sellerId > 'Chroma' and skuid > 'ChromaSKU1';
```

InvalidRequest: code=2200 [Invalid query] message="Only EQ and IN relation are supported on the partition key (unless you use the token() function)"

We can perform range operations efficiently only if the stored data is sorted

Partition key restrictions

Let's do a range query using ' $>$ ' on the partitioner key

```
cassandra@cqlsh:catalog> select * from skus where sellerId > 'Chroma' and skuid > 'ChromaSKU1';
```

```
InvalidRequest: code=2200 [Invalid query] message="Only EQ and IN relation are supported on the partition key (unless you use the token() function)"
```

Since, the partitions are **distributed** across the cluster, we cannot perform range operations on them

Partition key restrictions

Let's do a range query using ' $>$ ' on the partitioner key

```
cassandra@cqlsh:catalog> select * from skus where sellerId > 'Chroma' and skuid > 'ChromaSKU1';
```

InvalidRequest: code=2200 [Invalid query] message='Only EQ and IN relation are supported on the partition key (unless you use the token() function)'

We cannot use the $>$, $>=$, $<=$, $<$ operator directly on the partition key

Partition key restrictions

All the columns of the partition key
should be restricted in the query

We cannot use $>$, $>=$, $<=$, $<$ operator directly on the
partition key

Only IN and = operators are allowed on the
partition key

Partition key restrictions

Let's do a range query using '>' on the partitioner key

```
cassandra@cqlsh:catalog> select * from skus where sellerId > 'Chroma' and skuid > 'ChromaSKU1';
```

output

```
InvalidRequest: code=2200 [Invalid query] message="Only EQ and IN relation are supported on the partition key (unless you use the token() function)"
```

The coordinator node knows the token range of all nodes

Partition key restrictions

Range query using '>' using token (partition key)

```
cassandra@cqlsh:catalog> select * from skus where token(sellerid,skuid) >= token('Chroma', 'ChromaSKU1') ;
```

sellerid	skuid	timeofskucreation	productid	islistingcreated	listingid
title					
Fab	FABSKU1	2015-12-11 12:21	S0FA2	False	
Urban Living 2 seater					
Next	SKU1	2015-08-12 12:21	COM1	True	LISTINGNex
Acer One					
Fab	FABSKU	2015-12-11 12:21	S0FA1	True	LISTINGFab
Urban Living Derby					
Fab	FABSKU	2015-12-11 12:21	S0FA3	False	
Urban Living 4 seater					

(4 rows)

Partition key restrictions

All the columns of the partition key
should be restricted in the query

We cannot use $>$, $>=$, $<=$, $<$ operator directly on the
partition key

Only IN and = operators are allowed on the
partition key

We can use $>$, $>=$, $<=$, $<$ operator on the token

Partition key restrictions

Let's order the data by one of the partition key

```
cassandra@cqlsh:catalog> select * from skus where sellerid = 'Fab' AND skuid in ('FABSKU') order by sellerid;
```

InvalidRequest: code=2200 [Invalid query] message="Order by is currently only supported on the clustered columns of the PRIMARY KEY, got sellerid"

data cannot be ordered on partition key

Partition key restrictions

All the columns of the partition key
should be restricted in the query

We cannot use $>$, $>=$, $<=$, $<$ operator directly on the
partition key

Only IN and = operators are allowed on the
partition key

We can use $>$, $>=$, $<=$, $<$ operator on the token

ORDER BY is not supported with partition key

Let's query the skus CF with a restricted partition key

```
cassandra@cqlsh:catalog> select * from skus where sellerid in ('Decor', 'Fab') and  
sku_id IN('FABSKU', 'DECORSKU');
```

sellerid	sku_id	timeofskucreation	productid	islistingcreated	listingid
	title				
Decor	DECORSKU	2016-07-01 22:30	SOFA1	False	
null	Urban Living Derby				
Fab	FABSKU	2015-12-11 12:21	SOFA1	True	LISTINGFa
bSOFA1	Urban Living Derby				

(2 rows)

Partition key restrictions

All the columns of the partition key should be restricted in the query

We cannot use $>$, $>=$, $<=$, $<$ operator directly on the partition key

Only IN and = operators are allowed on the partition key

We can use $>$, $>=$, $<=$, $<$ operator on the token

ORDER BY is not supported with partition key