

Example 11:

cqlsh

Restrictions on clustering
keys

Let's see the skus CF description

```
cassandra@cqlsh:catalog> describe skus;
```

```
CREATE TABLE catalog.skus (  
    sellerid text,  
    skuid text,  
    timeofskucreation text,  
    productid text,  
    islistingcreated boolean,  
    listingid text,  
    title text,  
    PRIMARY KEY ((sellerid, skuid), timeofskucreation, productid)  
) WITH CLUSTERING ORDER BY (timeofskucreation ASC, productid ASC)  
    AND bloom_filter_fp_chance = 0.01  
    AND caching = {'keys': 'ALL', 'rows_per_partition': 'NONE'}  
    AND comment = ''  
    AND compaction = {'class': 'org.apache.cassandra.db.compaction.SizeTieredCompactionStrategy',  
'max_threshold': '32', 'min_threshold': '4'}  
    AND compression = {'chunk_length_in_kb': '64', 'class':  
'org.apache.cassandra.io.compress.LZ4Compressor'}  
    AND crc_check_chance = 1.0  
    AND dclocal_read_repair_chance = 0.1  
    AND default_time_to_live = 0  
    AND gc_grace_seconds = 864000  
    AND max_index_interval = 2048  
    AND memtable_flush_period_in_ms = 0  
    AND min_index_interval = 128  
    AND read_repair_chance = 0.0  
    AND speculative_retry = '99PERCENTILE';
```

Let's see the skus CF description

```
cassandra@cqlsh:catalog> describe skus;
```

```
CREATE TABLE catalog.skus (  
  sellerid text,  
  skuid text,  
  timeofskucreation text,  
  productid text,  
  islistingcreated boolean,  
  listingid text,  
  title text,
```

**data within a partition key will be
ordered first by timeofskucreation
and then by productid in ASC order**

```
PRIMARY KEY ((sellerid, skuid), timeofskucreation, productid)
```

) WITH CLUSTERING ORDER BY (timeofskucreation ASC, productid ASC)

```
AND bloom_filter_fp_chance = 0.01
```

```
AND caching = {'keys': 'ALL', 'rows_per_partition': 'NONE'}
```

```
AND comment = ''
```

```
AND compaction = {'class': 'org.apache.cassandra.db.compaction.SizeTieredCompactionStrategy',  
'max_threshold': '32', 'min_threshold': '4'}
```

```
AND compression = {'chunk_length_in_kb': '64', 'class':  
'org.apache.cassandra.io.compress.LZ4Compressor'}
```

```
AND crc_check_chance = 1.0
```

```
AND dclocal_read_repair_chance = 0.1
```

```
AND default_time_to_live = 0
```

```
AND gc_grace_seconds = 864000
```

```
AND max_index_interval = 2048
```

```
AND memtable_flush_period_in_ms = 0
```

```
AND min_index_interval = 128
```

```
AND read_repair_chance = 0.0
```

```
AND speculative_retry = '99PERCENTILE';
```

Let's query skus for Fab seller and sku FABSKU

```
cassandra@cqlsh:catalog> select * from skus where sellerid = 'Fab' and  
skuId = 'FABSKU';
```

sellerid	skuId	timeofskucreation	productid	islistingcreated	listingid	
	title					
	Fab	FABSKU	2015-12-11 12:21	SOFA1	True	LISTINGFabS
SOFA1	Urban Living Derby					
	Fab	FABSKU	2015-12-11 12:21	SOFA3	False	
null	Urban Living 4 seater					

timeofskucreation is same for both the rows
the row with SOFA1 comes before row with SOFA3

Let's query skus for Fab seller and sku FABSKU

```
cassandra@cqlsh:catalog> select * from skus where sellerid = 'Fab' and  
skuid = 'FABSKU';
```

	sellerid	skuid	timeofskucreation	productid	islistingcreated	listingid
	title					
	Fab	FABSKU	2015-12-11 12:21	S0FA1	True	LISTINGFabS
OFA1	Urban Living Derby					
	Fab	FABSKU	2015-12-11 12:21	S0FA3	False	
null	Urban Living 4 seater					

A more accurate way to represent how data is stored is
a **nested** structure

```
{
  'timeofskucreation': '2015-12-11 12:21'
    { 'productid': 'SOFA1'
      {
        'islistingcreated': 'True',
        'listingid': 'LISTINGFabSOFA1'
        'title': 'Urban Living Derby'
      }
      'productid': 'SOFA3'
      {
        'islistingcreated': 'False',
        'title': 'Urban Living 4 Seater'
      }
    }
}
```

```
'timeofskucreation': '2016-01-01 00:01'
```

```
. . .
```


**1st clustering column is the
outermost column of the
nested structure**

```
{  
  'timeofskucreation': '2015-10-01 00:01',  
  { 'productid': 'SOFA1',  
    { 'islistingcreated': 'True',  
      'listingid': 'LS100',  
      'title': 'Urban Living Derby'  
    }  
    , 'productid': 'SOFA3',  
    {  
      'islistingcreated': 'False',  
      'title': 'Urban Living 4 Seater'  
    }  
  }  
}
```

...

```
'timeofskucreation': '2016-01-01 00:01'
```

```
{
  'timeofskucreation': '2015-12-11 12:21'
  { 'productid': 'SOFA1'
    {
      'islistingcreated': 'True',
      'listingid': 'LISTINGFabSOFA1',
      'title': 'Urban Living 4 Seater'
    }
    'productid': 'SOFA3'
    {
      'islistingcreated': 'False',
      'title': 'Urban Living 4 Seater'
    }
  }
}
```

**timeofskucreation is the
first clustering column**

```
'timeofskucreation': '2016-01-01 00:01'
```

. . .


```
{  
  'timeofskucreation': '2015-12-11 12:21'  
    { 'productid': 'SOFA1'  
      {  
        'islistingcreated': 'True',  
        'listingid': 'LISTINGfabSOFA1',  
        'title': 'Urban living sofa'  
      }  
    }  
    { 'productid': 'SOFA3'  
      {  
        'islistingcreated': 'False',  
        'title': 'Urban living 4 seater'  
      }  
    }  
  }  
}
```

Next layer is the
2nd clustering
column ie productid

```
'timeofskucreation': '2016-01-01 00:01'
```

. . .

```
{  
  'timeofskucreation': '2015-12-11 12:21'  
  { 'productid': 'SOFA1'  
    {  
      'islistingcreated': 'True',  
      'listingid': 'LISTINGFabSOFA1',  
      'title': 'Urban Living Derby'  
    }  
  }  
  { 'productid': 'SOFA3'  
    {  
      'islistingcreated': 'False',  
      'title': 'Urban Living 4 Seater'  
    }  
  }  
}
```

All the remaining
columns are stored in the
layer below productid

```
'timeofskucreation': '2016-01-01 00:01'  
...  
'timeofskucreation': '2016-01-01 00:01'
```

In order to retrieve data efficiently, we need
to restrict the clustering columns in the
order in which they are defined

To prevent inefficient queries,
cassandra doesn't allow out of order
restrictions on clustering columns

Let's understand this with an example

Let's put a restriction on the productid column

WITH CLUSTERING ORDER BY (timeofskucreation ASC,
productid ASC)

```
cassandra@cqlsh:catalog> select * from skus where sellerid = 'Fab' and skuid in  
( 'FABSKU', 'FABSKU1') and productid = 'S0FA1';
```

InvalidRequest: code=2200 [Invalid query] message="PRIMARY KEY column "productid" cannot be restricted as preceding column "timeofskucreation" is not restricted"

The query fails as we haven't
restricted timeofskucreation column

Clustering column restrictions

If you want to restrict the n th clustering column, you need to restrict all the preceding 1 to $n-1$ clustering columns

For the query to work we need to restrict timeofskucreation

```
cassandra@cqlsh:catalog> select * from skus where sellerid =  
'Fab' and skuid in ('FABSKU', 'FABSKU1') and  
timeofskucreation = '2015-12-11 12:21' and productid =  
'S0FA1';
```

sellerid	skuid	timeofskucreation	productid	islistingcreated	listingid
title					
-----+-----+-----+-----+-----+-----					
-----+-----					
Fab	FABSKU	2015-12-11 12:21	S0FA1	True	LISTINGFabS
OFA1	Urban Living Derby				

Restricting clustering columns with IN

```
cassandra@cqlsh:catalog> select * from skus where sellerid in ('Decor', 'Fab') and  
skuid IN('FABSKU', 'DECORSKU') AND timeofskucreation IN ('2015-12-11 12:21', '2016-  
07-01 22:30') and productid IN ('SOFA1', 'SOFA2');
```

sellerid	skuid	timeofskucreation	productid	islistingcreated	listingid
	title				
Decor	DECORSKU	2016-07-01 22:30	SOFA1	False	
null	Urban Living Derby				
Fab	FABSKU	2015-12-11 12:21	SOFA1	True	LISTINGFa
bSOFA1	Urban Living Derby				

(2 rows)

Similarly IN query works if both the clustering columns are restricted

Clustering column restrictions

If you want to restrict the n th clustering column, you need to restrict all the preceding 1 to $n-1$ clustering columns

IN and = operators can be used with all the clustering columns

Let's restrict the clustering columns with range operators

```
cassandra@cqlsh:catalog> select * from skus where sellerid in ('Decor', 'Fab') and  
skuid IN('FABSKU', 'DECORSKU') AND timeofskucreation >= '2015-08-01' and timeofskuc  
reation <= '2016-07-01 22:30';
```

sellerid	skuid	timeofskucreation	productid	islistingcreated	listingid
	title				
Decor	DECORSKU	2016-07-01 22:30	SOFA1	False	
null	Urban Living Derby				
Fab	FABSKU	2015-12-11 12:21	SOFA1	True	LISTINGFa
bSOFA1	Urban Living Derby				

(2 rows)

Here productid is not restricted, which works perfectly fine
as data is sorted by timeofskucreation

Let's restrict the clustering columns with range operators

```
cassandra@cqlsh:catalog> select * from skus where sellerid in ('Decor', 'Fab') and  
skuid IN('FABSKU', 'DECORSKU') AND timeofskucreation >= '2015-08-01' and timeofskuc  
reation <= '2016-07-01 22:30' and productid IN ('SOFA1');
```

InvalidRequest: code=2200 [Invalid query] message="Clustering column "productid" cannot be restricted (preceding column "timeofskucreation" is restricted by a non-EQ relation)"

This query could result in scanning of entire data on a single node

But it will lead to read timeouts if the partitions are on different nodes

Clustering column restrictions

If you want to restrict n th clustering column, you need to restrict all the preceding 1 to $n-1$ clustering columns

IN and = operators can be used with all the clustering columns

Clustering column cannot be restricted if its preceding column is restricted by a range operator